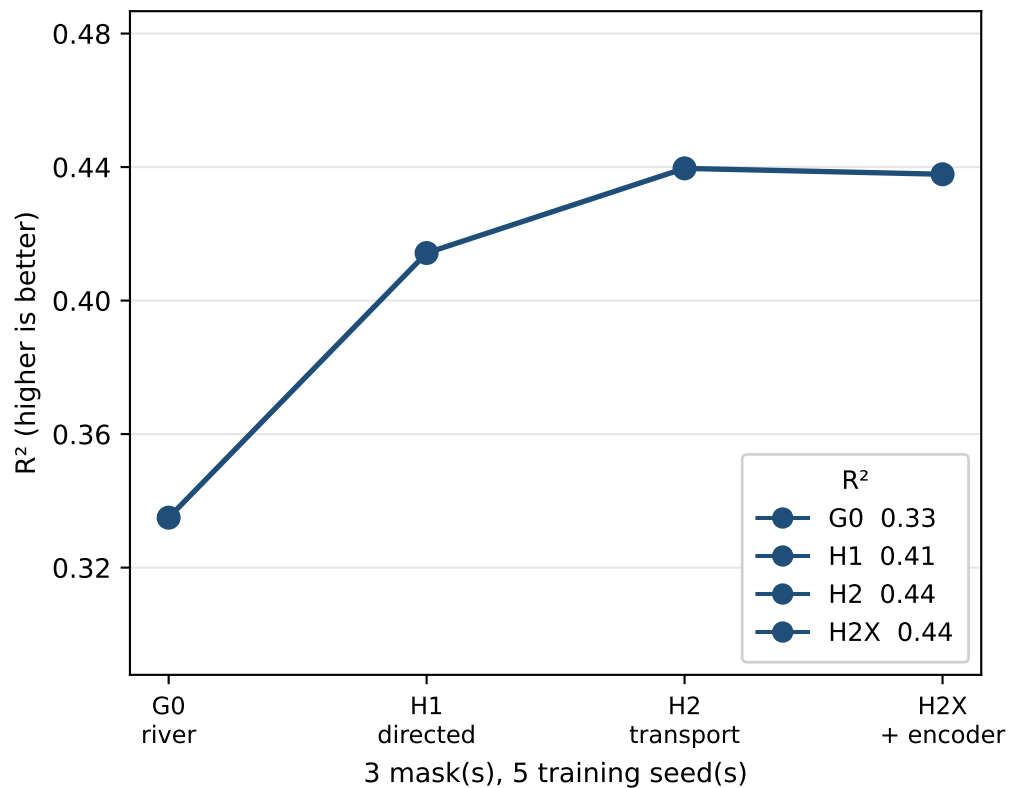
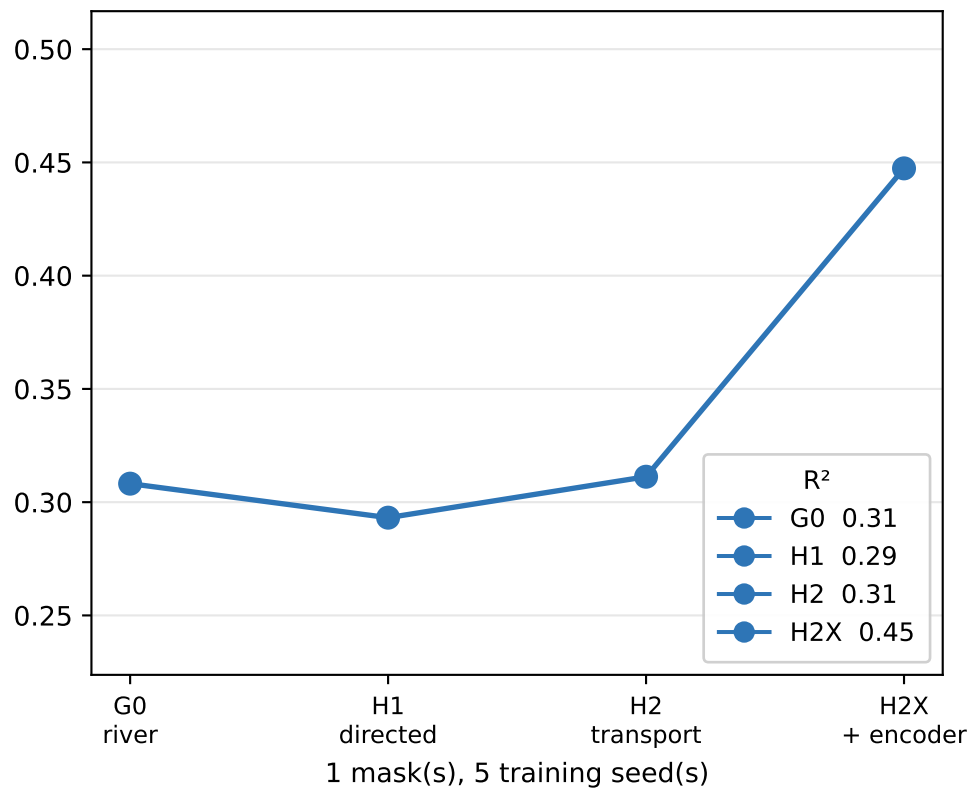


Figure 1 | Stored model performance across scenarios (R^2 , equal-weight mean over each scenario's masks)

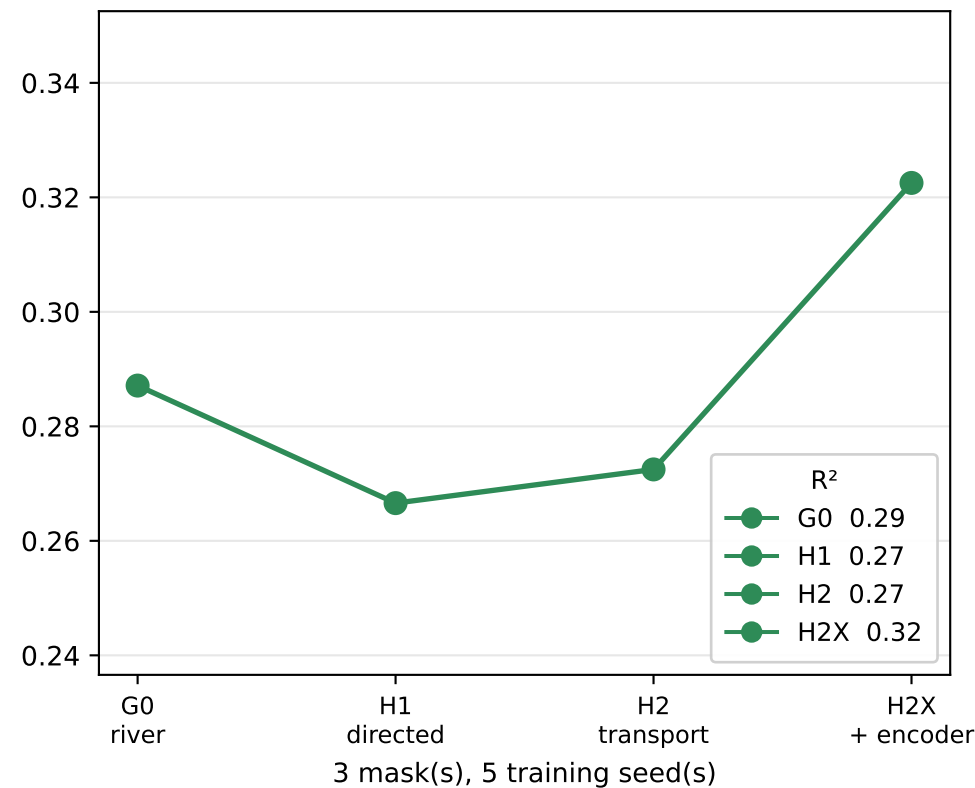
G0 plain GCN on the river graph
E1 random mask (r20)



H1 + upstream/downstream direction
E2b future reconstruction



H2 + transport gates
H2X + ecological context encoder
E3 unseen stations



G0: frozen single run. H1/H2/H2X: mean over 5 training seeds. A mean of per-mask R^2 is not the R^2 of all cells pooled. No error bars are shown because G0 has no training-seed spread to compare against. H2X adds ecological information AND an encoder at once, so the H2X gain is not attributable to the encoder alone. This is not a strict single-factor ablation.