



Figure 1 | Stored model performance across scenarios (R^2 , equal-weight mean over each scenario's masks)

G0 plain GCN on the river graph
E1 random mask (r20)




3 mask(s), 5 training seed(s)

H1 + upstream/downstream direction
E2b future reconstruction



1 mask(s), 5 training seed(s)

H2 + transport gates
H2X + ecological context encoder



3 mask(s), 5 training seed(s)

G0: frozen single run. H1/H2/H2X: mean over 5 training seeds. A mean of per-mask R^2 is not the R^2 of all cells pooled. No error bars are shown because G0 has no training-seed spread to compare against. H2X adds ecological information AND an encoder at once, so the H2X gain is not attributable to the encoder alone. This is not a strict single-factor ablation.